






Gaussian policy entropy (global)

Actual RL learner logs; not independent F/S success
Policy KL (miniepoth average) $1e-5$ LR applied on last optimizer update

Actor loss



Critic loss



Training episode reward (objectives differ)

