


Actual RL learner logs; not independent F/S success

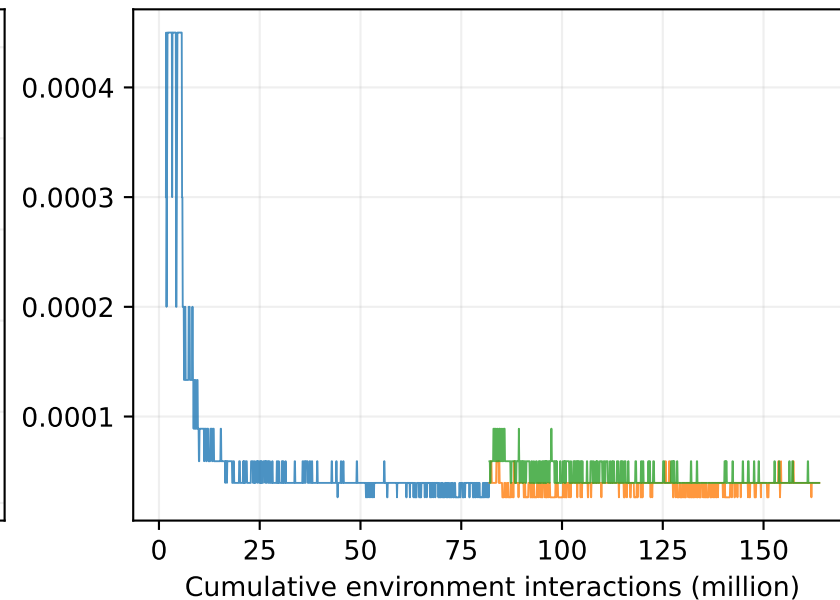
Gaussian policy entropy (global)



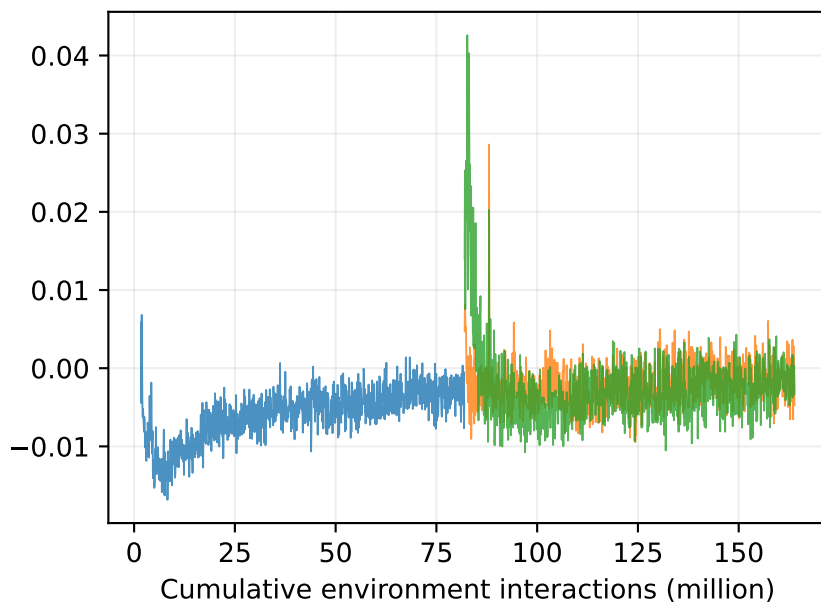
Policy KL (miniePOCH average)



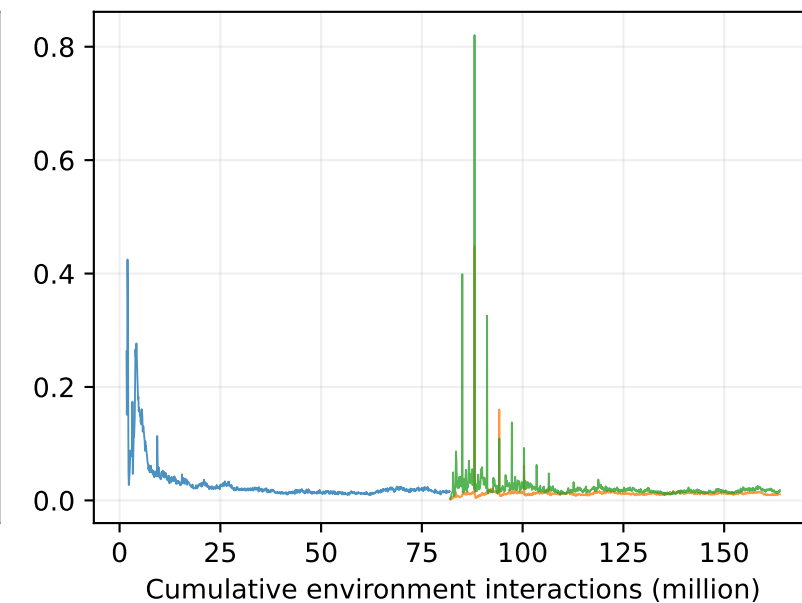
LR applied on last optimizer update



Actor loss



Critic loss



Training episode reward (objectives differ)

